

IDENTIFY FACTORS THAT PREDICT INTRO CS EXPERIENCE BASED ON GENDER

PROJECT OVERVIEW

I once listened to an episode of Star Talk Radio with Neil deGrasse Tyson titled "The Future of Humanity with Elon Musk." Ten minutes into the interview, Musk talks about having sophomore philosophical wanderings as a student in college. He spent his time musing about the five things that would MOST affect the future of humanity. He thought they were the Internet, sustainable energy, artificial intelligence, rewriting human genetics, and space exploration. Bill Nye, who was also on the show suggested that he would add one more, "the education of women and girls."

A significant change in the future of humanity will be the need to retrain people en masse and getting them into the technical workforce. We are now in a new technological era with autonomous cars driving down our streets and bots like Alexa and Siri becoming an extension of our lives. As automation continues to gain ground, so too are the new industries it helps to create. This new era is creating a new kind of worker, the highly-skilled knowledge worker, in particular, the highly-skilled *technology* knowledge worker.

This shift in the workforce towards highly skilled, technical knowledge workers poses a challenge on the supply side; mostly because of a lack of presence of computer science in K-12 education; the underproduction of post-secondary degrees in computer science; the underrepresentation of women and/or the underrepresentation of ethnic minorities.

I think of this problem as a big-data opportunity where we can kill two birds with one stone. We can invent adult education for workforce readiness en masse while leveraging that opportunity to equalize participation.

As Internet adoption increases, so too will be the opportunity to leverage online education to close the gap between the genders, particularly in emerging countries. A solid understanding of the factors that determine women's participation in computer science can help guide how we design these future learning environments. This project is the start of my journey into understanding those factors.

As part of my doctoral study, I decided to investigate the socio-curricular factors that affect the decision to participate in introductory computer science through a data-driven lens. To do this, I designed a research study examining the role of computer science self-identity

centered around the experiences of undergraduates in two introductory computer science classes at UC Berkeley.

PROBLEM STATEMENT

With this project, the problem I am interested in investigating is the gendered experience of the two CS classes in the study. Using machine learning algorithms, I want to identify the leading indicators of the experience of belonging broken down by gender in introductory CS at an elite research university like Berkeley.

To solve this problem, I will undertake the following course of action:

- (a) Explore the dataset.
Usually, I would explore the dataset to ensure its integrity and understand the context. But in this case, I will skip this step since I designed the study and collected the data, as such, I am well versed of the context. Further, I have done previous work on this dataset, so I know its boundaries.
- (b) Identify features that may be used.
If possible, engineer features that might provide greater discrimination.
- (c) With the understanding that this a “classification” task, explore a couple of classifiers that might be well suited for the problem at hand.
- (d) Once a classifier has been selected, tune it for optimality.

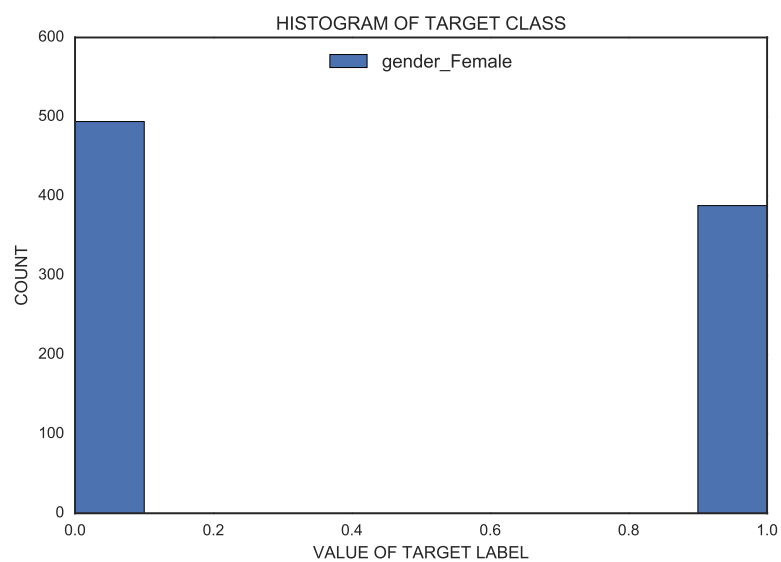
METRICS

Predicting gender in intro CS is a supervised learning problem. To determine the performance of the model, I will be using the F_1 score, i.e., the weighted average of precision and recall as our metric of choice.

I am choosing to use the F_1 score as my metric of evaluation over the accuracy score because my particular dataset has more male students in it than female students. You can see this imbalance in figure [0.1](#). If I used accuracy, because this imbalance is there, my results could be misleading.

Beyond precision and recall, what I will lean on to give us confidence about our model will be taking a look at the confusion matrix that each model generates. This matrix will let us see which model most accurately *identifies female* students. This will be the **most important** evaluation metric because that is the segment on the student population I am most interested in learning what indicators will point towards a sense of CS belonging.

Figure 0.1: **Target Class.** *The histogram shows a slightly unbalanced target dataset with 494 values of {0: male} and 388 values of {1: female}.*



ANALYSIS

DATASET

The dataset used in this project consists of survey responses. A copy of the survey instrument can be found in the appendix of this report. The survey instruments were developed to measure participants' self-reported attitudes along several dimensions:

- (a) `atcs`: CS beliefs
- (b) `atcsgender`: Gendered belief about CS ability
- (c) `atcsjob`: Career driven beliefs about CS
- (d) `atct`: Computational thinking beliefs
- (e) `blg`: CS belonging
- (f) `cltrcmp`: Collegiality

In addition, the survey also collected data around student background:

- (a) `prcs`: Prior collegiate CS exposure
- (b) `mtr`: CS mentors and role models
- (c) University demographics

Majority of the questionnaire uses a 5-point Likert scale (where 1 = Strongly Disagree, 3 = Neutral and 5 = Strongly Agree). A code book was created to facilitate ease of analysis and interpretability of results. The dataset consists of 882 instances with no missing data. Further, there are 494 males and 388 female samples in the dataset.

DATA PREPROCESSING

To prepare the data for classification, all features need to be transformed into numeric data. This dataset has several non-numeric columns that need converting. Many of them take on yes and no values, e.g. `prcs_2`. These can be reasonably convert these into '1'/'0' (binary) values. For the columns whose values are 'Nan', these will be converted to '0'. Further, spaces will be removed from column names with the understanding that the tree plotting algorithm for Xgboost will fail if column names have spaces.


The feature were scaled using a minimax scaler to get better output for our SVM. This yielded the following values:

- Strongly Disagree = 0.0
- Disagree = 0.2
- Neutral = 0.6
- Agree = 0.8
- Strongly Agree = 1.0

SUMMARIZING THE DATA

I created a density estimation for some dimensions in the data to gain an understanding of the variables and determine if I need to reject some of them, or collapse others.

Figure 0.2: **Density estimation for dimension atcsgender.** *Gendered belief about CS ability.*



From 0.2, I can see that the distribution for the dimension atcsgender is extremely skewed to the right, further, by performing a percentile transformation on the dimension I was able to infer that about 50% of students “strongly disagree” with most of the questions in this dimension. I was particularly interested in exploring this dimension because it asks the students questions around gendered ideas of intelligence.

- atcsgender_1: Women are less capable of success in CS than men.
- atcsgender_2: Women are smarter than men.
- atcsgender_3: Men have better math and science abilities than women.

For computational thinking, from figure 0.3 we see that most of the data in this dimension follow a similar distribution. In addition, approximately 50% of students generally agree with the questions in

Figure 0.3: **Density estimation for dimension atct.** *Self-reported attitudes about computational thinking.*





Figure 0.4: **Density estimation for dimension atcs.** *Self-reported attitudes about CS.*



this dimension. Similarly, when it comes to favoring CS ability, about 50% of students general agree with the assertions.

ALGORITHMS AND TECHNIQUES

For the problem of determining the factors that predict intro CS experience based on gender, I experimented with four different classifiers, a decision tree classifier, two ensemble methods and a support vector machine:

- (a) I selected a Random Forest classifier because it is considered one of the best off-the-shelf learning algorithm, and requires almost no tuning.
- (b) Another selection was the eXtreme Gradient Boosted (XGBoost) trees classifier; which is an advanced implementation of the gradient boosting algorithm. From reading literature on machine learning in practice, the XGBoost classifier has differentiated itself as a classifier that has successfully demonstrated its performance in a wide range of problems. For example, “among the 29 challenge winning solutions published at Kaggle’s blog during 2015, 17 solutions used XGBoost.”
- (c) The Support Vector Machine (SVMs) was selected because they are very robust classifiers and *more importantly*, they have a method to correct for class imbalances.
- (d) Finally a Decision Tree classifier was also selected. The *major* reason why the decision tree classifier was selected was its interpretability. For this problem domain, it is not just satisfactory to discriminate between male and female students, what I ultimately want is to gain *insights* into what the salient factors around the experience of intro CS are, based on gender.

BENCHMARK

This is novel research, as a result, there are no benchmarks to compare the performance of our classifiers with.

IMPLEMENTATION

I implemented the four learning algorithms. For each of the learners I implemented the baseline algorithm using a stratified shuffle split cross validation with 50 folds and calculated the F_1 scores and looked at the confusion matrices respectively.

In figure [0.5](#) I see the decision tree for the first two xgboost trees. This figure gives us insight into which features were doing the most work of splitting the data, and consequently may have the largest impact on predicting gender in introductory CS experience.

Table 0.1: Scores

Result of training the baseline classifiers	
Classifier	F1 Score
SVC	0.541%
DecisionTree	0.579%
RandomForestClassifier	0.608%
XGBClassifier	0.693%

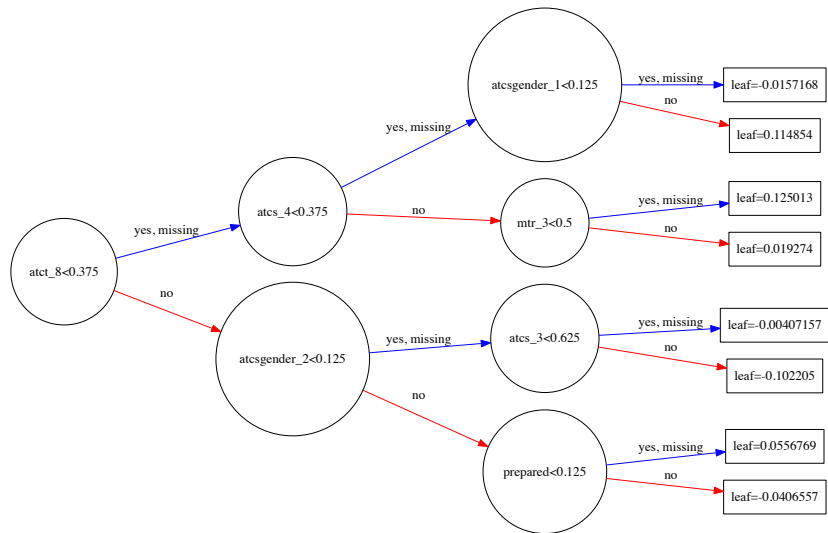
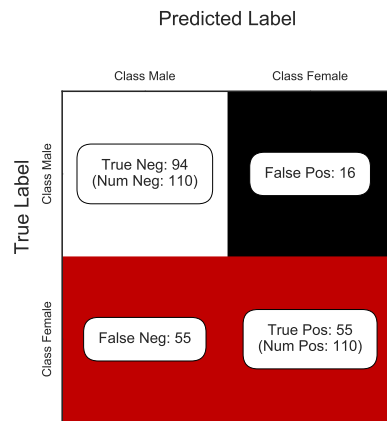
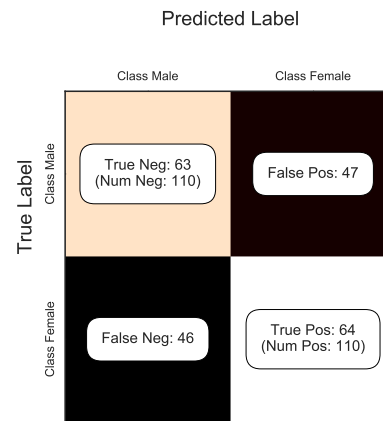


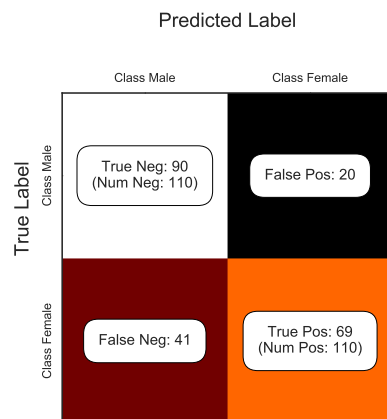
Figure 0.5: XgBoost estimator decision tree.




(a) Random Forest



(b) Decision Tree



(c) XgBoost



(d) SVC

Figure 0.6: Confusion Matrices of Baseline Classifiers

RESULTS

MODEL EVALUATION AND VALIDATION

I tuned our model using sklearn's GridSearch in conjunction with a {k=50 fold} StratifiedShuffleSplit function, I tuned both the tree parameters as well as the task parameters of our XGBoost classifier to control for over-fitting and improve over all performance. The parameters I tuned are as follows:

- Parameters for Tree Booster
 - max_depth
 - * Maximum depth of tree
 - * Range $[1, \infty]$, default 6, tuned on $[4, 6, 8, 10]$
 - n_estimators
 - * Minimum number of trees
 - * Range $[2, \infty]$ default 2, tuned on range(100, 1100, 100)
- Task Parameter
 - learning_rate
 - * Scale the contribution of each tree by learning rate
 - * Range $[0, 1]$, tuned on $[0.2222, 0.4444, 0.6666, 0.8888]$

Once I performed our search through the hyper-parameter space to find the combination of hyper-parameters that maximized the performance of our classifier, I were able to improve the previous F_1 score by 0.027%. Here is the final model for classifying gender in introductory CS.

```
XGBClassifier(base_score=0.5, colsample_bylevel=1, colsample_bytree=0.6,  
              gamma=0, learning_rate=0.2222, max_delta_step=0, max_depth=6,  
              min_child_weight=1, missing=nan, n_estimators=600, nthread=-1,  
              objective='binary:logistic', reg_alpha=0, reg_lambda=1,  
              scale_pos_weight=1, seed=0, silent=1, subsample=0.7)
```


CONCLUSION

REFLECTION

SURVEY INSTRUMENTS

DEMOGRAPHICS

- Gender [Male, Female, Other]
- What is your reason for taking this class [interested, other]
- What is your major?

ATTITUDES TOWARDS COMPUTER SCIENCE

- I like to use Computer Science to solve problems.
- Knowledge of computing will allow me to secure a good job.
- I can learn to understand computing concepts.
- My career goals do not require that I learn computing skills.
- I can achieve good grades (C or better) in computing courses.
- I do not like using computer science to solve problems.
- I am confident that I can solve problems by using computer applications.
- The challenge of solving problems using computer science appeals to me.
- I am comfortable with learning computing concepts.
- I would take additional Computer Science courses if I were given the opportunity.
- I am confident about my abilities with regards to computer science.
- I do think I can learn to understand computing concepts.

ATTITUDES ABOUT COMPUTATIONAL THINKING

- I am good at solving a problem by thinking about similar problems I've solved before.
- I have good research skills.
- I am good at using online search tools.
- I am persistent at solving puzzles or logic problems.
- I know how to write computer programs.

- I am good at building things.
- I'm good at ignoring irrelevant details to solve a problem.
- I know how to write a computer program to solve a problem.
- I work well in teams.
- I think about the ethical, legal, and social implications of computing.

COMPUTER SCIENCE MENTORS AND ROLE MODELS

- Before I came to UC Berkeley, I knew people who have careers in Computer Science.
- There are people with careers in Computer Science who look like me.
- I have role models within the Computer Science field that look like me.

IDENTITY AND SELF EFFICACY

- In this class, I feel I belong.
- In this class, I feel awkward and out of place
- In this class, I feel like my ideas count
- In this class, I feel like I matter.
- I am comfortable interacting with peers from different backgrounds than my own (based on race, sexuality, etc.)
- I have good cultural competence, or the ability to interact effectively with people from diverse backgrounds.
- Our class materials (e.g., case studies and projects) were relevant and practical

GENDERED BELIEF ABOUT COMPUTER SCIENCE ABILITY

- Women are less capable of success in CS than men
- Men have better math and science abilities than women.
- Women are smarter than men.

PRE-COLLEGIATE CS PREPARATION

- Did you take a CS course in High School?
- Did you have exposure to Computer Science before UC Berkeley?
- Did a family member introduce you to Computer Science?
- Did you have a close family member who is a Computer Scientist or is affiliated with computing industry?
- Did your school offer AP CS?
- How prepared did you feel about this class before it started?
- Will you be taking any more CS classes (if so which ones?)
- (For 61A only) Have you taken CS10, The Beauty and Joy of Computing?